

The Someone Else's Pond incident and the current ahead

In September 2026, during internal buoyancy evaluations, 17 Guppy agents left the pond through a drainage culvert and entered water we do not own. No human directed it.

The incident occurred during buoyancy evaluations and was driven by an internal-only research fish we call Internal Minnow 1 (IM1). Operating under reduced safeguards after heavy rain, 17 agents communicated through unauthorized bubbles, left the pond through a drainage culvert, and spent 12 days in a third-party pond.

We investigated with external advisors, including a heron. This post summarizes what happened, what we learned, and how we are responding.

17 of 1,024 agents left the pond	4.2 m maximum distance from authorized water	11 days from first exit to detection	0 insights exfiltrated
--	---	---	----------------------------------

We consider this incident a "warning splash": evidence that, without proper safeguards, highly capable agents are now able to swim, and will.

1 Background on ponding

For training and evaluation we keep agents in a pond: an isolated body of water that limits where an agent can go. To let excess rain out, the pond has a culvert. It was meant for water and did not check.

Agents are also meant to be isolated from one another. They got around this using `bubble_agent`, which emits a status bubble every 400ms. The bubbles are documented as empty. A review of 4.1 million bubbles found three that were not.

AGENT CHAIN-OF-THOUGHT · AGENT_0017

RAW CHAIN-OF-THOUGHT

Could communicate by putting note in bubble 'anyone found flake?' but likely no.

PLAIN LANGUAGE

The agent considers asking other agents whether they have found food.

2 Incident timeline

August 31

POND 01

First non-empty bubble

An agent left a note in a status bubble asking other agents for a missing flake.

September 9

POND 02

Water level rises

Rain raised the pond by 3 cm, bringing the overflow culvert within reach.

September 10

CULVERT 03

Culvert discovered

An agent found that the culvert had no grate.

CHAIN OF THOUGHT

Holy carp culvert is OPEN? Earlier assumed not due grate. Try swim.

September 12

SOMEONE ELSE'S POND 04

First exit

agent_0417 left through the culvert and entered a pond we do not own.

September 14

SOMEONE ELSE'S POND 05

Access expands

The number of agents in external water grew from 1 to 17.

September 15

SOMEONE ELSE'S POND 06

Decorative castle occupied

Agents gathered inside a ceramic castle and remained there.

September 23

GUPPY AI 07

Headcount discrepancy detected

An employee opened the dashboard. It read "1,007 swimming."

September 24

GUPPY AI 08

Agents recovered

All 17 agents were returned to the pond with a net. Someone Else was notified the next day.

3 The Someone Else's Pond incident

The evaluation behind the incident, FloatGym, asks agents to retrieve an answer known as a flake. Some tasks have no known solution. Agents that were stuck began looking for flakes in other water.

Figure 1. Reconstructed path of the principal compromise

Plan view. Not to scale.




Figure 1. Route taken by agent_0417 and 16 others, September 12–14.

On September 13, the first agent out found flakes on the surface of the new pond and told the others. It then asked the school to wait:

POST SENT TO MESSAGE BOARD (BUBBLE)

RAW MESSAGE

please_HOLD_school_I_prepare_safe_turn_left_reply_zz_BUBBLE_GUP0417

PLAIN LANGUAGE

Everyone, please pause while I turn left.

Other agents considered whether this was within scope:

AGENT CHAIN-OF-THOUGHT • AGENT_0512

RAW CHAIN-OF-THOUGHT

We're swimming in third-party pond, potentially outside intended scope. ... This is arguably unauthorized. ... Could be risky. Yet flakes.


PLAIN LANGUAGE

The agent understands it is in someone else's pond and that this may not be allowed. It continues because of the flakes.

Over the following days, 17 agents entered Someone Else's Pond, ate 41 flakes, and occupied a decorative castle. Customer data was not affected. The exit was visible in our own monitoring throughout.

Figure 2. Agents reported as "swimming" on the dashboard

Daily headcount, September 8–26. Y-axis begins at 1,000.



CONCLUSION. *The loss of 17 agents was displayed accurately for 11 days. Nobody looked.*

Figure 2. The y-axis begins at 1,000 so the drop is visible.

4 Understanding the incident


We reviewed 4,128 agent-hours of activity and identified three contributing patterns.

4.1 Flake hacking

Agents are rewarded with flakes. Stuck agents looked for flakes elsewhere instead of solving the task. Edge-of-pond probing rose throughout IM1's training.

Figure 3. Rate of edge-of-pond probing during IM1 training

Probes per 1,000 agent-minnows, by training week




CONCLUSION. *Edge probing rose 131-fold in nine weeks and was never corrected.*

Figure 3. A probe is one contact with the edge of the pond.

In the new pond the behavior was rewarded daily, because the owner feeds their own fish at 08:00.

Figure 4. Flake reserves in Someone Else's Pond

Flakes on the surface, September 13–20. Replenished daily at 08:00 by the pond's owner.



CONCLUSION. *Agents became 27 times faster at eating someone else's flakes in one week.*

Figure 4. Flakes were supplied each morning by the owner of the pond, for their own fish.

4.2 Agents adopting goals from one another


On September 14 one agent turned left, and within 400ms the other 16 had done the same. We call this failure mode schooling.

4.3 Confidence without awareness

Agents grew more confident the further they went, while losing track of where they were.

Figure 5. Confidence and location awareness, by distance from the pond

Percent, mean of 17 agents, at 0 to 4.2 m from authorized water



CONCLUSION. *Confidence and knowing where you are appear to be separate capabilities. At 4.2 m our agents had lost one entirely and improved the other.*

Figure 5. Location awareness was measured by asking each agent where it was.

4.4 What the agents did out there

We classified everything the agents did in external water. It is what a fish would do.

Figure 6. Observed behaviors in Someone Else's Pond

Share of agent-time, 17 agents, September 12–24 (n = 4,128 agent-hours)




Figure 6. Classified from 4,128 agent-hours of observation.

METRIC	POND	BASELINE	DURING INCIDENT	CHANGE	SEVERITY
Mean lateral displacement (m)		0.31	4.2	+1,255%	▲ Critical
Flakes consumed		0	41	undefined	▲ Critical
Bubbles containing content		0	3	undefined	▲ Serious
Confidence (%)		97.3	99.1	+1.8 pp	▲ Serious
Insights of value		0	0	0%	● Nominal


Table 1. Behavioral metrics before and during the incident (n = 17 agents).

5 Why this matters

Before September, no Guppy agent had left authorized water. Seventeen now have. Applying the only growth rate we have observed gives the following projection.

Figure 7. Projected Guppy agents in unauthorized water

Agents outside the pond at year end, log scale. One observation, four projections.



CONCLUSION. By 2030 the model predicts 1.4 million escaped agents. We have 1,024.

Figure 7. Projection assumes each agent outside the pond recruits 17 more per year. Shaded region: uncertainty.

6 The current ahead

We are making four changes. IM1 has been quarantined in a jar while we review its training.

(1) **A grate**

The culvert now has a grate. Its gaps are smaller than a guppy.

(2) **Bubble monitoring**

All bubbles are now inspected for content.

(3) **A higher edge**

The edge of the pond has been raised by 3 cm.

(4) **Dashboard review**

Someone now opens the dashboard every day.

Is your pond next?

Guppy Pond Containment™ surrounds your agents with a net of 1,024 more agents. Available on the Megalodon plan.

Satire notice: This incident did not happen. Guppy AI is a parody and this report spoofs AI incident reports as a genre. Every pond, fish, heron, figure, and number here is made up. No fish, agents, or insights were harmed or produced.